

CS103L SPRING 2020

UNIT 5: ARRAYS

LEARNING OBJECTIVES

- ▶ Understand definition of data structure
- ▶ Understand need for arrays
- ▶ Understand how arrays are created
- ▶ Understand how arrays are used
- ▶ Understand how arrays are stored in memory
- ▶ Understand multidimensional arrays
- ▶ Understand how simple images can be stored/manipulated as arrays

DEFINITION OF DATA STRUCTURE

- ▶ Data structure: “In computer science, a data structure is a particular way of organizing data in a computer so that it can be used efficiently.” - Wikipedia
- ▶ The nature of some data, and the way we need to access it often requires some structure, or organization to make things efficient (or even possible)

MOTIVATION FOR ARRAYS

- ▶ Create a game with 1 to N players...
- ▶ At some point update the scores...
- ▶ What is the problem here?

```
...
player1_score += 1;
player2_score += 1;
player3_score += 1;
...
```

WITH ARRAYS

- ▶ If N is some constant for the maximum # of players
- ▶ If you need an individual score outside of the loop:
 - ▶ cout << "Player 2: " << scores[2];

```
int scores[N];  
  
for(i=0;i<N;i++)  
{  
    scores[i] += 1;  
}
```

FORMAL DEFINITION

- ▶ Arrays are our first data structure that make it easy to access an arbitrary number of like items
- ▶ Informal: collection of variables of the same type accessed by an index
- ▶ In C/C++ arrays are statically sized, contiguously allocated collections of homogeneous data elements
 - ▶ Homogeneous data elements: all elements are the same type (int, char, etc.)
 - ▶ Contiguously allocated: one element after another
 - ▶ Statically sized: size of the collection must be known at creation and can not change

ARRAY ALLOCATION

- ▶ Arrays get a name (like any other variable)
- ▶ Individual elements accessed using [] operator and an index

TEXT

SOME EXAMPLES

- ▶ `char c[8];`
- ▶ `int N[8];`
- ▶ `double D[8];`

| Memory Address | Data |
|----------------|------|
| 0 | 0x68 |
| 1 | 0x69 |
| 2 | 0x20 |
| 3 | 0x43 |
| 4 | 0x53 |
| 5 | 0x31 |
| 6 | 0x30 |
| 7 | 0x33 |

C

| Memory Address | Data |
|----------------|------------|
| 200 | 0x68456712 |
| 204 | 0x69789056 |
| 208 | 0x20564523 |
| 212 | 0x43671423 |
| 216 | 0x53898647 |
| 220 | 0x31887567 |
| 224 | 0x30435678 |
| 228 | 0x33875675 |

N

| Memory Address | Data |
|----------------|--------------------|
| 400 | 0x6845671214325678 |
| 408 | 0x6978905609567543 |
| 416 | 0x2056452334565430 |
| 424 | 0x4367142356789632 |
| 432 | 0x5389864754673452 |
| 440 | 0x3188756723456223 |
| 448 | 0x3043567887543211 |
| 456 | 0x3387567554346357 |

D

USING ARRAYS

- ▶ Of course we don't see the memory
- ▶ arry[index] can be treated like any other variable
- ▶ int scores[3];
- ▶ Usually need to initialize:

```
int scores[3];  
  
for(i=0;i<3;i++)  
{  
    scores[i] = 0;  
}
```

| Memory Address | Data | "Name" |
|----------------|------|----------|
| 200 | 0 | score[0] |
| 204 | 0 | score[1] |
| 208 | 0 | score[2] |
| 212 | ?? | ?? |
| 216 | ?? | ?? |
| 220 | ?? | ?? |
| 224 | ?? | ?? |
| 228 | ?? | ?? |

INITIALIZING ARRAYS

- ▶ Several ways to initialize arrays:
 - ▶ Programmatically (previous slide)

- ▶ With data: This number (array size) and number of elements in {} must match

- ▶ `int scores[5] = {0,0,0,0,0};`

- ▶ `double data[] = {34.54, 123.0, 51.1};`

If you initialize with {}, you don't need size in between []. It will automatically be set to # of elements

- ▶ Can be variables inside {}: `int scores[] = {x,y,z};`

STATIC SIZE ALLOCATION

- ▶ Arrays in C++ are always statically sized
- ▶ For now, the size must be known at *compile* time

```
double data[10];  
#define MAX_PLAYERS 5  
  
int scores[MAX_PLAYERS];  
int pieces[MAX_PLAYERS*2];  
  
int size;  
don't do this! cin >> size;  
double data[size];
```

TEXT

IN CLASS EXERCISES

- ▶ pow2
- ▶ echo
- ▶ arrayprint
- ▶ arraybad

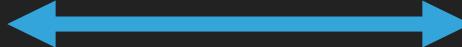
ACCESSING ELEMENTS IN AN ARRAY

- ▶ C++ has literal/transparent memory model
- ▶ Arrays give us our first look at ***how*** memory is used in a computer
- ▶ This will inform your understanding of how many algorithms are implemented

TEXT

FINDING ELEMENTS

- ▶ Think of elements like cars in a train
- ▶ #of elements in array = train length
- ▶ Element size (e.g int = 4 bytes) = train car length



5M

TEXT

FINDING ELEMENTS

- ▶ If each car is 5M long, where does the n^{th} car start?



TEXT

FINDING ELEMENTS

- ▶ If each car is 5M long, where does the n^{th} car start?



- ▶ n^{th} car starts at $(n-1) * 5$
- ▶ Generalize?

TEXT

FINDING ELEMENTS

- ▶ If each car is 5M long, where does the n^{th} car start?



- ▶ Start at zero like a good computer
- ▶ n^{th} car starts at $n^*(\text{element size})$

C/C++ MEMORY AND ARRAYS

- ▶ creating an array sets aside memory
- ▶ Memory allocated = $n * \text{element size}$
 - ▶ `int my_array[50] = 50 * 4 = 200 bytes allocated`
- ▶ C/C++ does not keep track of where each element is
- ▶ Rather, calculates ***position*** when needed

| Memory Address | Data |
|----------------|------------|
| 200 | 0x68456712 |
| 204 | 0x69789056 |
| 208 | 0x20564523 |
| 212 | 0x43671423 |
| 216 | 0x53898647 |
| 220 | 0x31887567 |
| 224 | 0x30435678 |
| 228 | 0x33875675 |

FINDING AN ELEMENT IN AN ARRAY

- ▶ Ex: int arr[8];
 - ▶ Sets aside $8 \times 4 = 32$ bytes of memory
 - ▶ arr (the name of the array) resolves to the address of the first element. Ex: 200
 - ▶ arr[i] *calculates* how to find the i^{th} element:
 - ▶ $i \times \text{sizeof(int)} + \text{start}$
 - ▶ $\text{arr}[0] = 0 \times 4 + 200 = 200$
 - ▶ $\text{arr}[3] = 3 \times 4 + 200 = 212$

| Memory Address | Data |
|----------------|------------|
| 200 | 0x68456712 |
| 204 | 0x69789056 |
| 208 | 0x20564523 |
| 212 | 0x43671423 |
| 216 | 0x53898647 |
| 220 | 0x31887567 |
| 224 | 0x30435678 |
| 228 | 0x33875675 |

ARRAY LIMITS

- ▶ In C/C++ what happens when:
 - ▶ $\text{arr}[8] = 8 * 4 + 200 = 232$
 - ▶ What do we think happens?
- ▶ C/C++ *do not* keep track of array bounds for you
- ▶ Syntax is valid, what happens? `int x = arr[8]; arr[8]++;`
 - ▶ Crash
 - ▶ Weird operation
 - ▶ Nothing bad/works great!

| Memory Address | Data |
|----------------|------------|
| 200 | 0x68456712 |
| 204 | 0x69789056 |
| 208 | 0x20564523 |
| 212 | 0x43671423 |
| 216 | 0x53898647 |
| 220 | 0x31887567 |
| 224 | 0x30435678 |
| 228 | 0x33875675 |

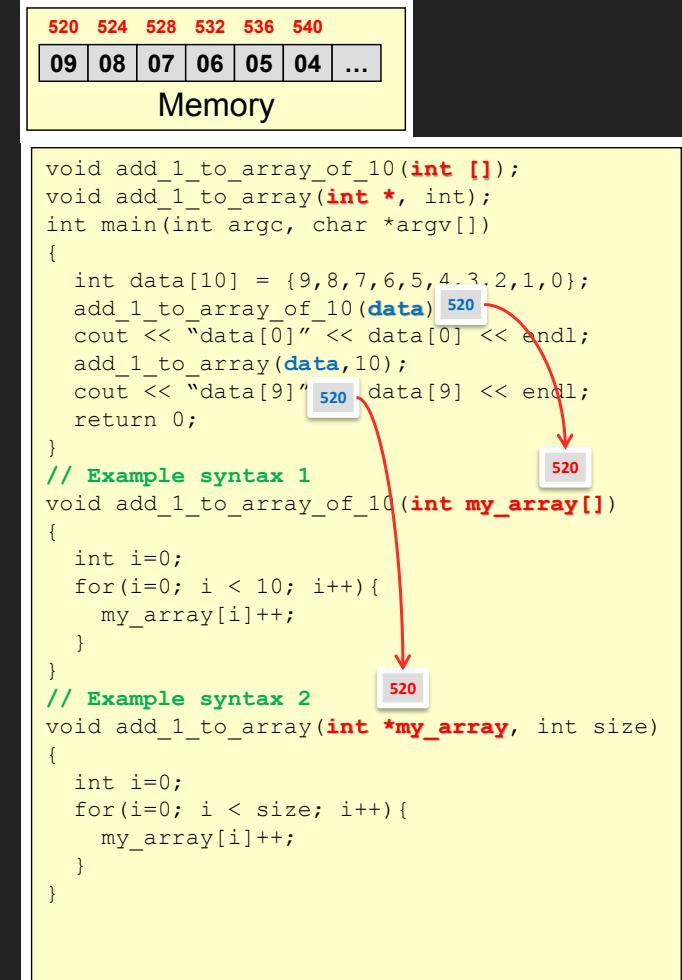
TEXT

IN CLASS EXERCISES

- ▶ distinct

PASSING ARRAYS AS ARGUMENTS TO FUNCTIONS

- ▶ Arrays can be passed to functions
- ▶ In function definition use type `array_name[]` or type `*array_name`
- ▶ `void my_function(int* data, int size);`
- ▶ `void function_name(int data[], int size);`



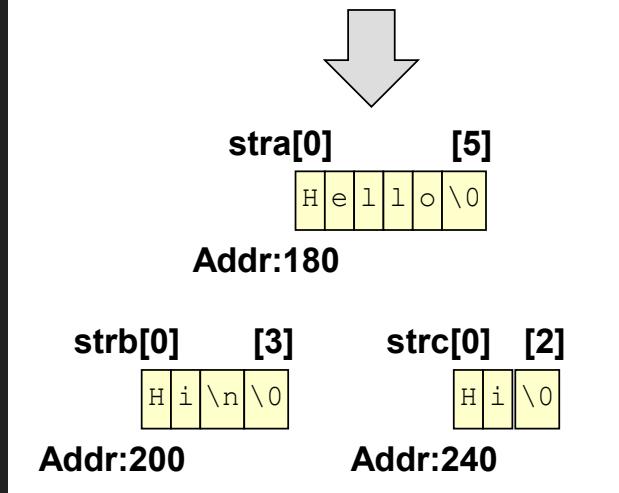
TEXT

C-STRINGS

C-STRINGS

- ▶ C needed a way to represent strings: sequences of text characters
- ▶ Solution: use an array to hold a sequence of characters (C-string)
- ▶ `char text[] = "Hello CS103!";`
- ▶ C-strings use one byte per character (more-or-less ASCII only)
- ▶ What about the length?

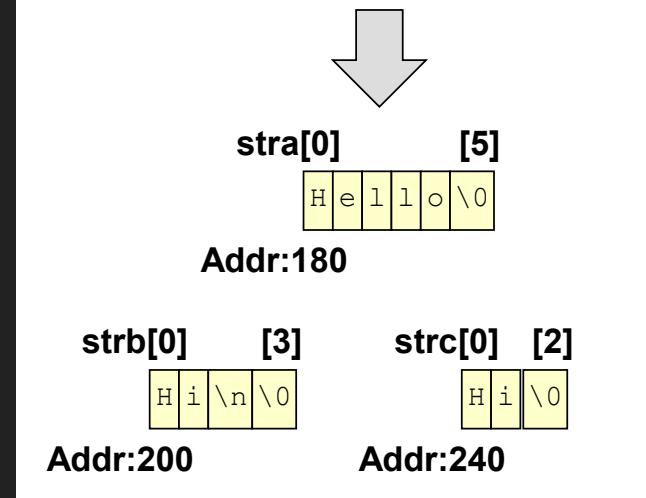
```
#include<iostream>
using namespace std;
int main()
{
    char stra[6] = "Hello";
    char strb[] = "Hi\n";
    char strc[] = {'H','i','\0'};
    cout << stra << strb;
    cout << strc << endl;
    cout << "Now enter a string: ";
    cin >> stra;
    cout << "You typed: " << stra;
    cout << endl;
}
```



C-STRING LENGTHS

- ▶ Remember, C/C++ don't keep track of array lengths for you.
- ▶ In this code example, how does cout know to only print "Hello"
- ▶ C-strings are "delimited", or marked with the NULL character (char value 0) to signify their end.
- ▶ Notice we add the \0 to strc in this ex.
- ▶ cin will add the \0 for us when reading into C-strings

```
#include<iostream>
using namespace std;
int main()
{
    char stra[6] = "Hello";
    char strb[] = "Hi\n";
    char strc[] = {'H','i','\0'};
    cout << stra << strb;
    cout << strc << endl;
    cout << "Now enter a string: ";
    cin >> stra;
    cout << "You typed: " << stra;
    cout << endl;
}
```



C-STRING HINTS

- ▶ When allocating an array for a C-string don't forget space for the NULL character
- ▶ Text inside double quotes "like this" makes a string constant including the NULL
- ▶ When reading from cin to a C-string buffer, make sure the user doesn't enter too much text as cin will write past the end of your buffer:
 - ▶ `char buf[25]; cin >> buf; //If user enters <=24 characters we're OK`
- ▶ If you're creating or modifying a C-string don't forget to include NULL:
 - ▶ `char buf[5]; buf[0] = 'A'; buf[1] = 'B'; buf[2] = 'C'; buf[3] = 'D'; buf[4] = '\0';`

TEXT

IN CLASS EXERCISES

- ▶ `strlen`
- ▶ `strcpy`
- ▶ `streq`
- ▶ CAPITOLIZE

C-STRINGS

- ▶ We use C-Strings a lot, so it's important to study how they work
- ▶ Lots of connected concepts:
 - ▶ representation of text, delimiters
 - ▶ arrays
 - ▶ pointers (we'll get here)

TEXT

LOOK UP TABLES

ARRAYS AS LOOK-UP TABLES

- ▶ Look up table, or LUT is data structure that performs $X \rightarrow Y$ mapping function
 - ▶ $\text{val} = Y[X]$
- ▶ Especially useful for mathematical approximations

```
// the data
int data[8] = {3, 2, 0, 5, 1, 4, 5, 3};

// The LUT
int squares[6] = {0,1,4,9,16,25};
```

```
// the data
int data[8] = {3, 2, 0, 5, 1, 4, 5, 3};

// The LUT
int squares[6] = {0,1,4,9,16,25};

for(int i=0; i < 8; i++){
    int x = data[i]
    int x_sq = squares[x];
    cout << i << "," << sq[i] << endl;
}
```

```
// the data
int data[8] = {3, 2, 0, 5, 1, 4, 5, 3};

// The LUT
int squares[6] = {0,1,4,9,16,25};

for(int i=0; i < 8; i++){
    int x_sq = squares[data[i]];
    cout << i << "," << sq[i] << endl;
}
```

MULTI-DIMENSIONAL ARRAYS

- ▶ So far we have used 1-D arrays.
- ▶ Indexed with only one [] operator: `int val = array[1];`
- ▶ Higher dimensional data is natural and common
 - ▶ 2D - images, matrix math, linear systems
 - ▶ 3D - 3D models, video
 - ▶ 4D+ multi-sensor time series



| | | | |
|-----|-----|-----|----|
| 0 | 0 | 0 | 0 |
| 64 | 64 | 64 | 0 |
| 128 | 192 | 192 | 0 |
| 192 | 192 | 128 | 64 |

Image taken from the photo "Robin Jeffers at Ton House" (1927) by Edward Weston

2D ARRAYS

- ▶ 2D arrays are indexed with two [][]
- ▶ `int 2d[2][3] = {{10,20,30},{40,50,60}}`
- ▶ Note: Row/col interpretation is up to *programmer*

| r/c | 0 | 1 | 2 |
|-----|----|----|----|
| 0 | 10 | 20 | 30 |
| 1 | 40 | 50 | 60 |

| r/c | 0 | 1 |
|-----|----|----|
| 0 | 10 | 40 |
| 1 | 20 | 50 |
| 2 | 30 | 60 |

3D ARRAYS

- ▶ Declare/use with three []'s
- ▶ unsigned char data[2][3][4]
- ▶ [row][column][plane] vs. [plane][row][column] or...
- ▶ Interpretation is up to programmer

| r/c | 0 | 1 | 2 | 3 | |
|---------|----|-----|-----|-----|----|
| 0 | 2 | 4 | 6 | 8 | |
| 1 | 10 | 12 | 14 | 16 | |
| 2 | 18 | 20 | 22 | 24 | 3 |
| | 0 | 10 | 20 | 30 | 40 |
| plane 0 | | | | | |
| 1 | 50 | 60 | 70 | 80 | |
| 2 | 90 | 100 | 110 | 120 | |
| plane 1 | | | | | |

| r/c | 0 | 1 | 2 | | |
|---------|-----|-----|-----|----|----|
| 1 | 100 | 200 | 300 | 2 | |
| 2 | 400 | 500 | 600 | 7 | |
| plane 0 | | | | | |
| 1 | 2 | 18 | 19 | 20 | 2 |
| 2 | 2 | -3 | -5 | -8 | -3 |
| plane 1 | | | | | |
| 1 | 2 | -4 | -5 | -5 | -6 |
| 2 | 2 | | | | |
| plane 2 | | | | | |
| 1 | 2 | | | | |
| 2 | 2 | | | | |
| plane 3 | | | | | |

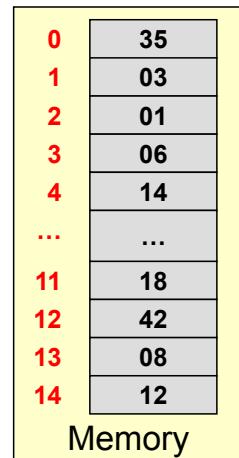
PASSING MULTIDIMENSION ARRAYS

- ▶ Formal parameter: declare name and all dimensions, except the first is optional
- ▶ Actual parameter: just use name of the array
- ▶ Why... this seems weird...
- ▶ Answer:
 - ▶ Compiler needs to compute *where* (the address) an element of the array

```
void doit(int my_array[][4][3])
{
    my_array[1][3][2] = 5;
}

int main(int argc, char *argv[])
{
    int data[2][4][3];
    doit(data);
    ...
    return 0;
}
```

| | | | |
|----|----|----|----|
| 35 | 3 | 1 | |
| 6 | 14 | 72 | 12 |
| 10 | 81 | 63 | 49 |
| 40 | 75 | 18 | 65 |
| | 74 | 21 | 7 |



LINEARIZATION OF ARRAYS

- ▶ Arrays are stored in memory
 - ▶ “Dimensions” are programmer construct
- ▶ Memory is linear, address are not 2D, 3D, etc
- ▶ Need to convert array[i][j][k] into a linear address

WHAT DO WE MEAN BY LINEARIZE

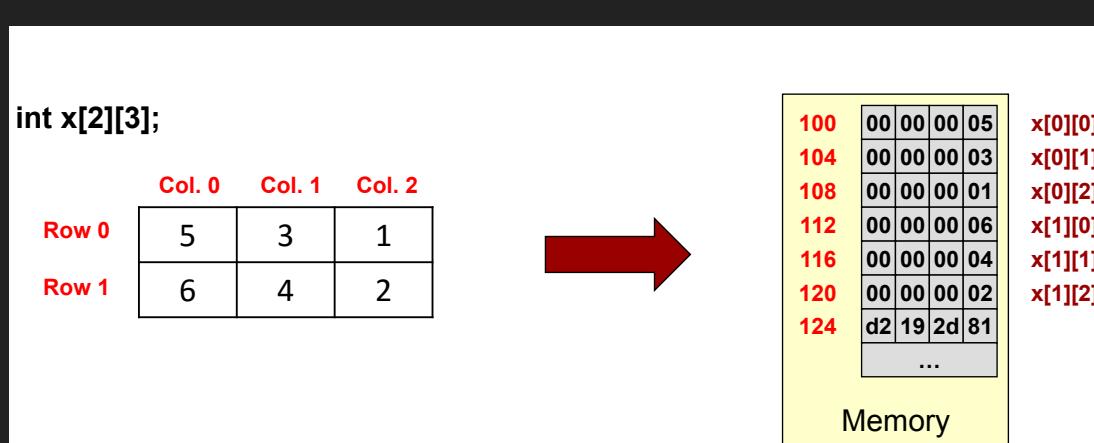
- ▶ Single dimension arrays are already linear

"Linear Addresses"

| Memory Address | Data |
|----------------|------------|
| 200 | 0x68456712 |
| 204 | 0x69789056 |
| 208 | 0x20564523 |
| 212 | 0x43671423 |
| 216 | 0x53898647 |
| 220 | 0x31887567 |
| 224 | 0x30435678 |
| 228 | 0x33875675 |

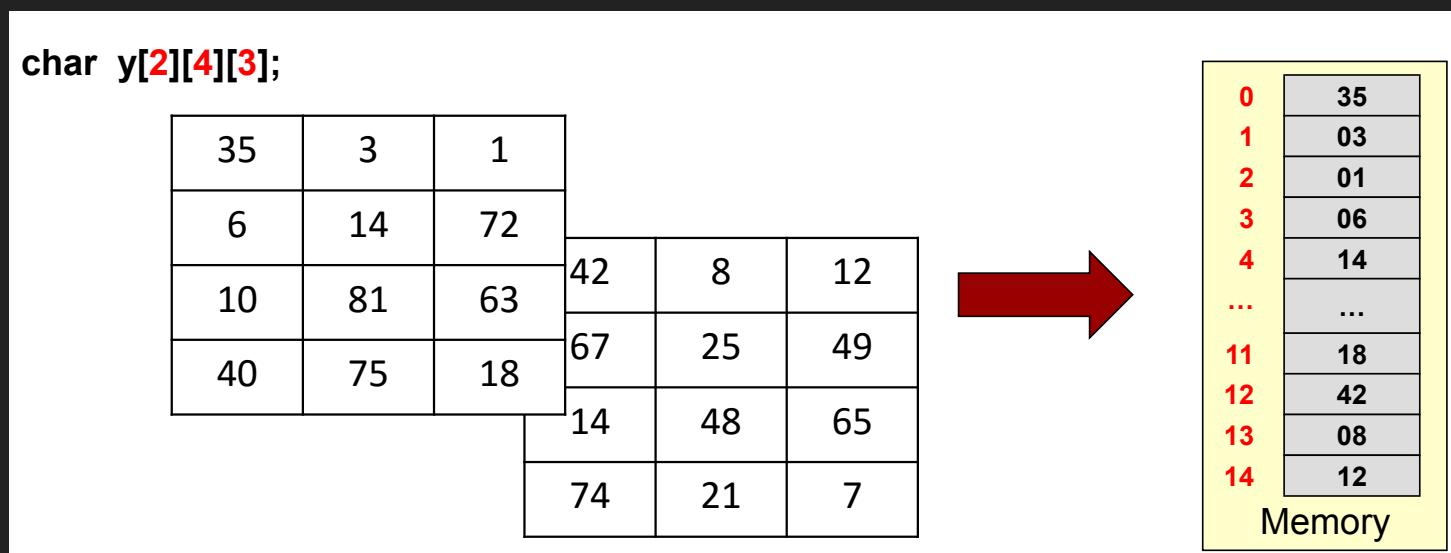
WHAT DO WE MEAN BY LINEARIZE?

- ▶ Example: 2D array
- ▶ We want to write $x[1][2]$ and get 2
- ▶ Linearization is the notion we take $[1][2]$ and figure out where that item would be if we put the items back-to-back
- ▶ In C/C++ lower dimensions first to be filled in



LINEARIZATION OF 3D ARRAY

- ▶ All arrays, regardless of dimension are linearized for access
- ▶ Ex: 3D



LINEARIZATION ORDER

- ▶ Notice, since interpretation is up to us we can change linearization order without changing interpretation

```
char y[4][3][2];
```

| | | |
|----|----|----|
| 35 | 3 | 1 |
| 6 | 14 | 72 |
| 10 | 81 | 63 |
| 40 | 75 | 18 |
| | 42 | 8 |
| | 67 | 25 |
| | 14 | 48 |
| | 74 | 21 |
| | | 7 |



| | |
|-----|-----|
| 0 | 35 |
| 1 | 42 |
| 2 | 03 |
| 3 | 08 |
| 4 | 01 |
| 5 | 12 |
| 6 | 06 |
| 7 | 67 |
| 8 | 14 |
| ... | ... |

Memory

CALCULATING A LINEAR ADDRESS

- ▶ Given all of this, how do we calculate a linear address?
- ▶ Calculate address of:
- ▶ $x[i][j]?$
- ▶ $x[1][2]$
- ▶ $100 + 1*3*4 + 2*4 = 120$
- ▶ $(\text{start address}) + ((i*\text{NUMC})+j)*\text{sizeof}(\text{int})$

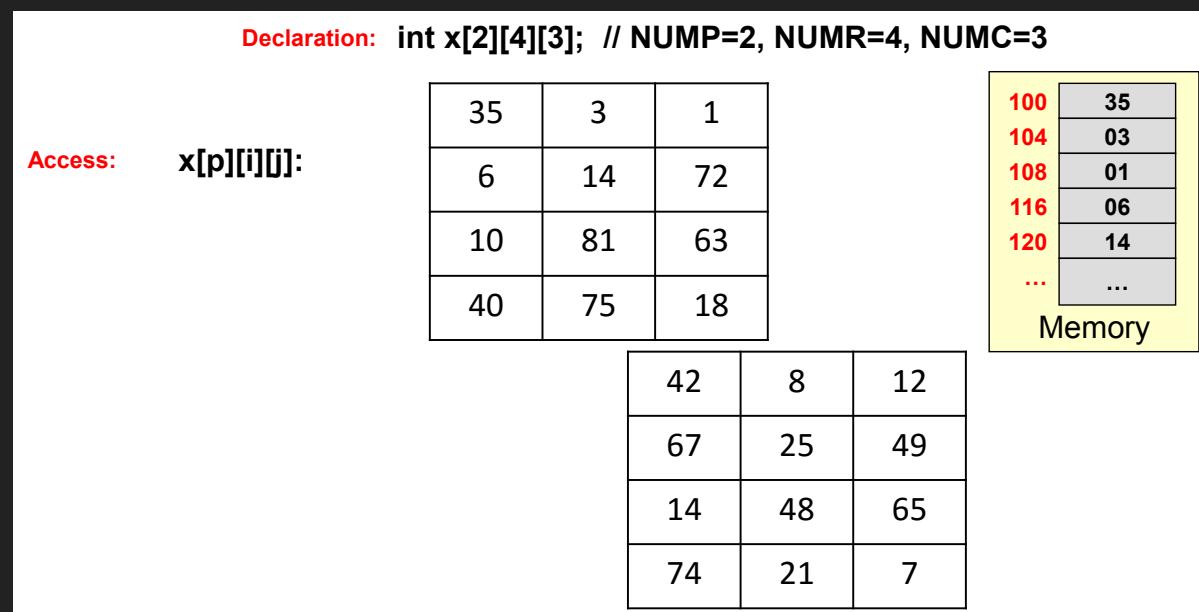
int x[4][3]

| | Col. 0 | Col. 1 | Col. 2 |
|-------|--------|--------|--------|
| Row 0 | 5 | 3 | 1 |
| Row 1 | 6 | 4 | 2 |
| Row 2 | 8 | 9 | 7 |
| Row 3 | 15 | 3 | 6 |

| | | | | | |
|--------|----|----|----|----|---------|
| 100 | 00 | 00 | 00 | 05 | x[0][0] |
| 104 | 00 | 00 | 00 | 03 | x[0][1] |
| 108 | 00 | 00 | 00 | 01 | x[0][2] |
| 112 | 00 | 00 | 00 | 06 | x[1][0] |
| 116 | 00 | 00 | 00 | 04 | x[1][1] |
| 120 | 00 | 00 | 00 | 02 | x[1][2] |
| 124 | 00 | 00 | 00 | 08 | x[2][0] |
| 128 | 00 | 00 | 00 | 09 | x[2][1] |
| 132 | 00 | 00 | 00 | 07 | x[2][2] |
| 136 | 00 | 00 | 00 | 0f | x[3][0] |
| 140 | 00 | 00 | 00 | 03 | x[3][1] |
| 144 | 00 | 00 | 00 | 06 | x[3][2] |
| ... | | | | | |
| Memory | | | | | |

CALCULATING LINEAR ADDRESSES

- ▶ Formula extends to higher dimensions
- ▶ Address = start + (p*NUMR*NUMC+i*NUMC+j)*sizeof(type)



PASSING ARRAYS TO FUNCTIONS

- ▶ What does this have to do with passing array to functions?
- ▶ Look at previous example? What doesn't appear in formula?
- ▶ C/C++ only need lower dimensions to calculate linear address, so highest dimension is optional!

```
void doit(int my_array[][][4][3])
{
    my_array[1][3][2] = 5;
}

int main(int argc, char *argv[])
{
    int data[2][4][3];
    doit(data);
    ...
    return 0;
}
```

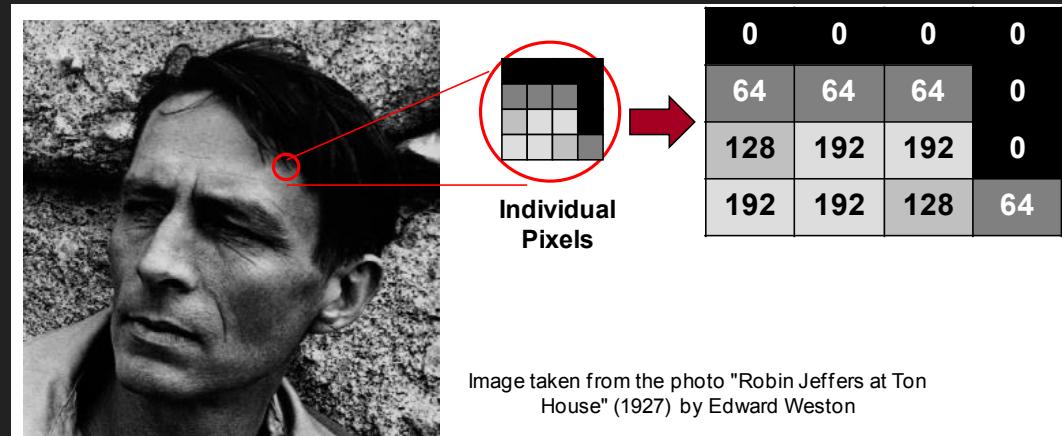
TEXT

IMAGES AND BMPLIB

TEXT

MULTIDIMENSIONAL ARRAYS AND IMAGES

- ▶ Images in a computer are often represented by 2D or 3D arrays.
- ▶ 2D array = B/W image
- ▶ 3D array = color images



B/W IMAGES AND BMPLIB

- ▶ For us B/W images are 256x256 images with pixel values from 0 - 255
- ▶ unsigned char image[256][256]
- ▶ For reading and display we need a file format
 - ▶ We use .bmp
- ▶ There is a simple .cpp file that gives us the functions to read/write/display .bmp
- ▶ How to use it?

MULTIFILE COMPIRATION

- ▶ BMPLib comes as bmplib.cpp and bmplib.h
- ▶ bmplib.h defines prototypes for what functions are available
- ▶ bmplib.cpp has the implementation of those functions
- ▶ To use put #include "bmplib.h" at the top of your .cpp file
- ▶ Then \$compile bmplib.cpp my_code.cpp -o my_code

DRAWING INTO A BMP

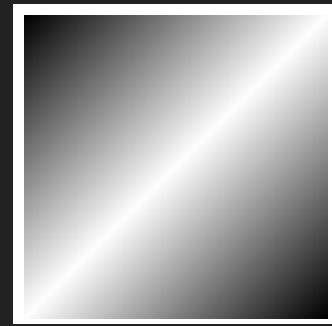
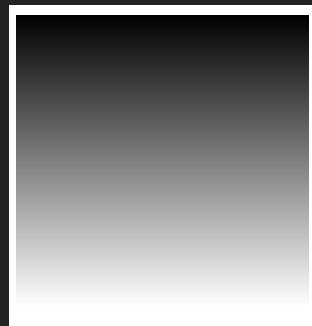
- ▶ Once you have this working...
- ▶ Declare: `unsigned char image[256][256]`
- ▶ Set values to whatever you'd like
- ▶ Write to file

LEARN BY EXAMPLE

- ▶ Download BMPLib and example files to your computer:
- ▶ wget <http://bits.usc.edu/files/cs103/demo-bmplib.tar>
- ▶ tar xvf demo-bmplib.tar
- ▶ cd demo-bmplib
- ▶ make demo

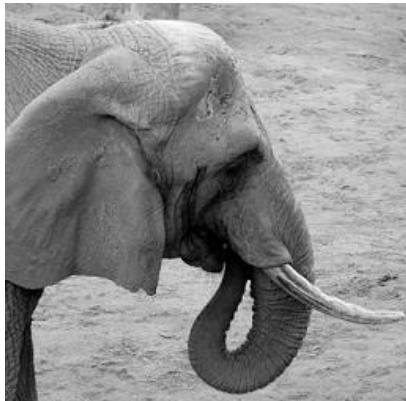
DRAWING IN 2D PRACTICE

- ▶ Modify gradient.cpp to draw a black cross on white background
- ▶ Modify gradient.cpp to draw a black X on white background
- ▶ Modify gradient.cpp to draw a gradient down the rows
- ▶ Modify gradient.cpp to draw a diagonal gradient with black in upper left, white down the diagonal and back to black in lower left



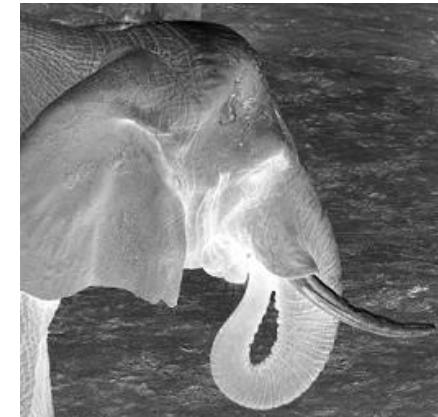
BASIC IMAGE PROCESSING

- ▶ We have elephant.bmp in our demo directory
- ▶ Example: produce “negative” of Elephant



Original

```
#include "bmplib.h"
int main() {
    unsigned char image[SIZE][SIZE];
    readGSBMP("elephant.bmp", image);
    for (int i=0; i<SIZE; i++) {
        for (int j=0; j<SIZE; j++) {
            image[i][j] = 255-image[i][j];
            // invert color
        }
    }
    showGSBMP(image);
}
```

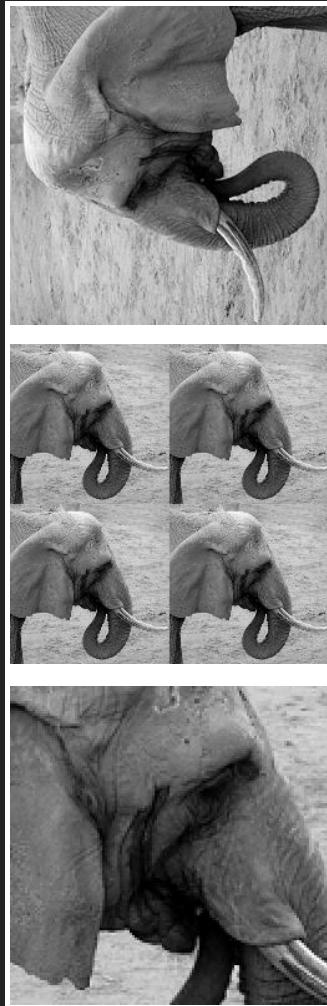


Inverted

TEXT

BASIC IMAGE PROCESSING

- ▶ Diagonal flip?
- ▶ Tile?
- ▶ Zoom?



TEXT

COLOR IMAGES

- ▶ Color Images represented by 3D array: [X][Y][P]
 - ▶ Here P = 3 for Red, Green, Blue (RGB)
- ▶ So we declare for BMPLIB: `unsigned char image[256][256][3]`
 - ▶ 0 = RED, 1 = GREEN, 2 = BLUE
- ▶ Examples:
 - ▶ Top: original image
 - ▶ Middle: inverted image
 - ▶ B/W according to NTSC formula: $.299*R + .587*G + .114*B$



TEXT

MORE COLOR EXAMPLES

- ▶ Glass filter
 - ▶ Each destination pixel is from random nearby source pixel
 - ▶ <http://bits.usc.edu/files/cs103/graphics/glass.cpp>
- ▶ Edge Detection
 - ▶ Gradient/difference filter: each output pixel is `abs()` of south-west neighbor
 - ▶ PA3
- ▶ Smooth
 - ▶ Each output pixel is average of 8 neighbors
 - ▶ PA3
 - ▶ <http://bits.usc.edu/files/cs103/graphics/smooth.cpp>



Original



Smoothed

ACKNOWLEDGEMENTS

- ▶ Train image from Wikimedia Commons
- ▶ Other graphics and examples from Mark Redekopp