

CS103L SPRING 2020

UNIT 7: FILE I/O

I/O STREAMS IN C++

- ▶ I/O: short for input/output
- ▶ Older term from mainframe days for getting data into and out of your program
- ▶ C++ offers object oriented abstractions to make using I/O easier
- ▶ We use the “stream” abstractions through cin, cout, ifstreams, ofstreams

STREAM ABSTRACTION

- ▶ The stream objects we use model I/O as a “stream”
- ▶ A stream is a series of data (usually single characters) one after the other
 - ▶ Narrow but long, like a stream
- ▶ We interact with streams differently depending on if they are input or output streams
 - ▶ Input streams, we extract characters from the stream into our program
 - ▶ Output streams, we insert characters from our program into the stream

INPUT STREAMS

- ▶ Input streams are a series of characters available for our program to read
- ▶ For now, just "cin" where the characters come from keyboard
- ▶ Next, file input streams
- ▶ Eventually internal input streams (stringstreams)
- ▶ Also, network streams

CIN INPUT STREAM

- ▶ cin keeps a buffer of characters read from the keyboard
- ▶ Allows us to read them one-by-one
- ▶ Or formatted - powerful abstraction!

```
char c;  
cin >> c;
```

H E L L O _ W O R L D ↵



- 3 4 . 1 ↵

```
double x;  
cin >> x;
```

COUT OUTPUT STREAM

- ▶ Using cout with the << operator takes data from our program and inserts it into an output stream
 - ▶ The stream is usually shown on the terminal
- ▶ Output is formatted: cout decides how to present the data
 - ▶ int or double: as a number
 - ▶ char: as a single character
 - ▶ char*: as a string (follow the pointer, print the chars until NULL)
 - ▶ Any other pointer type: memory address as hexadecimal

I/O STREAM ABSTRACTION

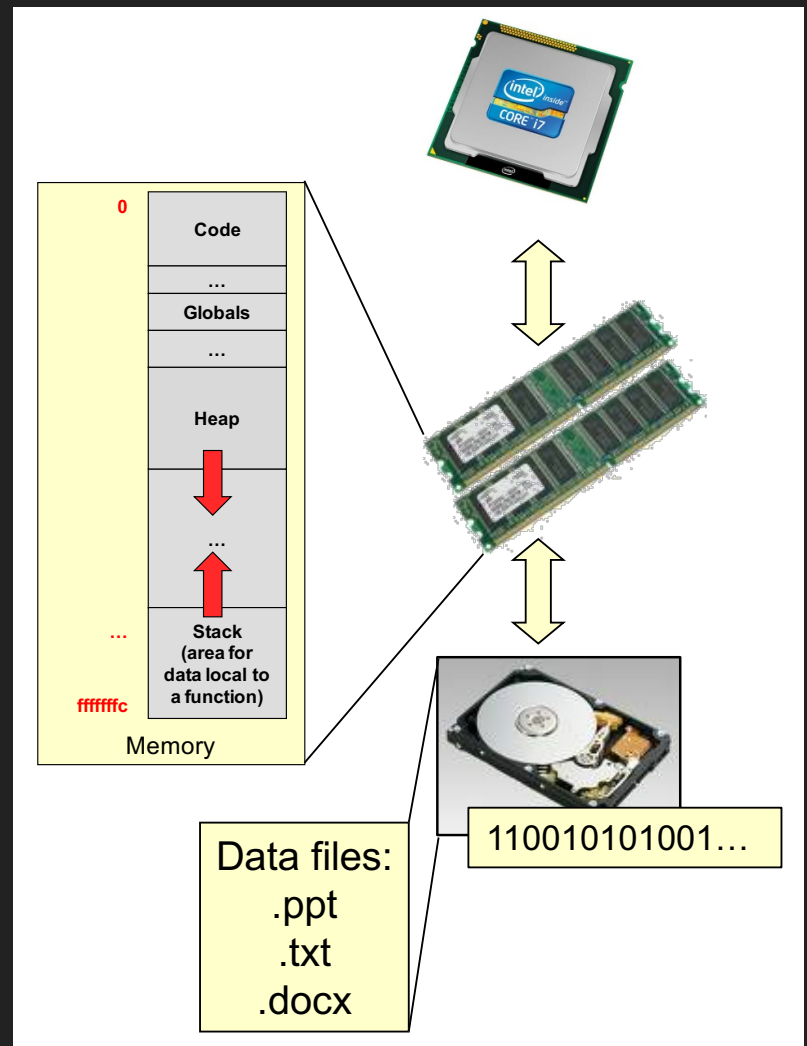
- ▶ Why is the I/O stream abstraction powerful?
- ▶ Once you can use one type of I/O stream, you can use any!
- ▶ They all define the same interface

FILE I/O

- ▶ How can we get data into our programs?
 - ▶ cin
 - ▶ Command line arguments
- ▶ How can we get data out?
 - ▶ Just cout for now
- ▶ Now we add file I/O through file streams

NOTE ON COMPUTER ARCHITECTURE

- ▶ CPU can talk to RAM and I/O devices
 - ▶ Like keyboard, terminal, disk controllers
- ▶ Our program can only interact with data in RAM, specifically *its* memory space only
- ▶ OS and C++ provide interfaces and abstractions so that C++ programs can do I/O (keyboard, terminal, disk, network)



IMPORTANT NOTE!

- ▶ Remember: everything we do in C++ is an interaction with a variable or memory address
- ▶ In order to work with data in C++ we must put it in a variable
 - ▶ After the data is loaded in a variable (or memory) we can process it
- ▶ Already seen with cin/cout, now adding file I/O
- ▶ Everything we're learning today is how to get data out of a file into a variable (and the other way around)

FILE I/O INTERFACE

- ▶ `#include <fstream>`
- ▶ `ifstream` object type
 - ▶ Lets us open a file, extract data, close file
- ▶ `ofstream` object type
 - ▶ Lets us open a file for writing, insert data, save and close the file

BINARY VS. TEXT DATA

- ▶ Granularity of data in file world is the byte
 - ▶ Files are modeled as streams of bytes, one-after-another
- ▶ Two types:
 - ▶ Text files: bytes are interpreted as a series of ASCII character
 - ▶ .txt file, code, csv file, HTML
 - ▶ Binary files: bytes are interpreted as binary data as specified by the file format:
 - ▶ Image files, videos, audio, compiled executables
- ▶ Example: store integer 103 in a file
 - ▶ Text file: '1' '0' '3' - takes up 3 bytes
 - ▶ Binary: could store as single byte 0x67 or if an 32-bit int, 0x00000067 - depends on file format

TEXT FILE I/O

- ▶ We only do text file I/O in 103
- ▶ `#include <fstream>`
- ▶ Use `ifstream` object to read from a file
 - ▶ Works **exactly** like `cin`, use `>>` operator
- ▶ Use `ofstream` object to write to a file
 - ▶ Works **exactly** like `cout` use `<<` operator

input.txt

5 -3.5

output.txt

Int from file is 5
Double from file is -3.5

```
#include <iostream>
#include <fstream>
using namespace std;

int main ()
{
    int x; double y;
    ifstream ifile ("input.txt");
    if( ifile.fail() ){ // able to open file?
        cout << "Couldn't open file" << endl;
        return 1;
    }

    ifile >> x >> y;
    if ( ifile.fail() ){
        cout << "Didn't enter an int and double";
        return 1;
    }

    ofstream ofile("output.txt");

    ofile << "Int from file is " << x << endl;
    ofile << "Double from file is " << y << endl;

    ifile.close();
    ofile.close();

    return 0;
}
```

GETTING LINES OF TEXT

- ▶ `>>` skips then stops at whitespace
- ▶ Sometimes we want a whole **line** of text
- ▶ `stream.getline(char *buf, int bufsize)`
 - ▶ `cin.getline(buf, 100); ifile.getline(buf, 100);`
- ▶ Program reads whole lines of text, prints out

```
#include <iostream>
#include <fstream>
using namespace std;


int main ()
{
    char myline[100]; int i = 1;
    ifstream ifile ("input.txt");
    if( ifile.fail() ){ // can't open?
        return 1;
    }

    ifile.getline(myline, 100);
    while ( ! ifile.fail() ) {
        cout << i++ << ": " << myline << endl;
        ifile.getline(myline, 100);
    }

    ifile.close();
    return 0;
}
```

input.txt

```
The fox jumped over the log.\n
The bear ate some honey.\n
The CS student solved a hard problem.\n
```



```
1: The fox jumped over the log.
2: The bear ate some honey.
3: The CS student solved a hard problem.
```

IN-CLASS EXERCISES

- ▶ wget http://ee.usc.edu/~redekopp/cs103/file_io_ex.tar
- ▶ sum_from_file
- ▶ reverse_it
- ▶ countnum

OTHER INPUT STREAM FEATURES

- ▶ Getting lines of text
 - ▶ Seen already with files, applies to cin as well
 - ▶ >> operator skips, then stops at whitespace
- ▶ Use stream.getline(char* buf, int max) to read a line of text including whitespace, but not the '\n')
- ▶ .getline() reads up to max-1 characters and inserts the NULL-char for you

```
#include <iostream>
using namespace std;

int main ()
{
    char mytext[80];
    cout << "Enter your full name" << endl;
    cin.getline(mytext, 80);

    int last=0;
    for(int i=0; i<80; i++){
        if(mytext[i] == ' '){
            last = i;
            break;
        }
    }
    cout << "Last name starts at index: ";
    cout << last << endl;
    return 0;
}
```


CHECKING INPUT STREAM ERRORS

- ▶ Until now we didn't know how to detect errors with cin (or ifstreams)
- ▶ Input streams have .fail() method - returns 'true' if something has gone wrong
 - ▶ Couldn't open file, asked for int but got 'abc', reached end of input file, etc.
- ▶ Once an error occurs .fail() returns 'true' until .clear() is called.

```
#include <iostream>
using namespace std;

int main ()
{
    int x;
    cout << "Enter an int: " << endl;
    cin >> x; // What if the user enters:
              //      "abc"

    // Check if we successfully read an int
    if( cin.fail() ) {
        cout << "Error: I said enter an int!";
        cout << " Now I must exit!" << endl;
        return 1;
    }

    cout << "You did it! You entered an int";
    cout << " with value: " << x;

    return 0;
}
```

TEXT

EXAMPLE

- ▶ `wget http://ee.usc.edu/~redekopp/cs103/cinfail.cpp`
- ▶ cinfail example

UNDERSTANDING INPUT STREAMS AND .FAIL()

. User enters value "512" at 1st prompt, enters "123" at 2nd prompt

```
int x=0;
```

X = 0 cin =

```
cout << "Enter X: ";
```

X = 0 cin = 5 1 2 \n

```
cin >> x;
```

X = 512 cin = \n

cin.fail() is false

```
int y = 0;
```

Y = 0 cin = \n

```
cout << "Enter Y: ";
```

Y = 0 cin = \n 1 2 3 \n

```
cin >> y;
```

Y = 123 cin = \n

cin.fail() is false

UNDERSTANDING INPUT STREAMS AND .FAIL()

. User enters value “23abc” at 1st prompt, 2nd prompt fails

```
int x=0;
```

X = cin =

```
cout << "Enter X: ";
```

X = cin =

```
cin >> x;
```

X = cin =

cin.fail() is **false**

```
int y = 0;
```

Y = cin =

```
cout << "Enter Y: ";
```

Y = cin =

```
cin >> y;
```

Y = cin =

cin.fail() is **true**

UNDERSTANDING INPUT STREAMS AND .FAIL()

. User enters value "23 99" at 1st prompt, 2nd prompt skipped

```
int x=0;
```

X = cin =

```
cout << "Enter X: ";
```

X = cin =

```
cin >> x;
```

X = cin =

cin.fail() is **false**

```
int y = 0;
```

Y = cin =

```
cout << "Enter Y: ";
```

Y = cin =

```
cin >> y;
```

Y = cin =

cin.fail() is **false**

UNDERSTANDING INPUT STREAMS AND .FAIL()

. User enters value “23 99” at 1st prompt, everything read as string

```
char x[80];
```

```
cout << “Enter X: “;
```

```
cin.getline(x, 80);
```

X = cin =

X = cin =

2	3		9	9	\n
---	---	--	---	---	----

X =

23	99
----	----

 cin =

cin.fail() is
false

**NOTE: \n character is
discarded!**

USING .FAIL, .IGNORE, .CLEAR

- ▶ You use `.fail()` to detect errors when reading data.
- ▶ `.ignore(int n, char delim)` clears out the buffer up to the next `delim` character, or a maximum of `n`
- ▶ `.clear()` resets `.fail()` flag, otherwise `.fail()` will continue to return 'true'

```
#include <iostream>
using namespace std;

int main ()
{
    int x;
    cout << "Enter an int: " << endl;
    cin >> x; // What if the user enters:
              //      "abc"

    // Check if we successfully read an int
    while( cin.fail() ) {
        cin.clear(); // turn off fail flag
        cin.ignore(256, '\n'); // clear inputs
        cout << "I said enter an int: ";
        cin >> x;
    }

    cout << "You did it!  You entered an int";
    cout << " with value: " << x;

    return 1;
}
```

OS REDIRECTION AND PIPES

- ▶ Redirection and pipes give a secondary way to do file I/O with programs
- ▶ They are OS level services
- ▶ Performed on the command line with:
 - ▶ '`<`' input redirection
 - ▶ '`>`' output redirection
 - ▶ `|` pipe output to input

INPUT REDIRECTION

- ▶ '`<`' input redirection
 - ▶ Takes the contents of a file and makes this the 'stdin' of your program
 - ▶ Extracting from cin with `>>` gets data from the file, not the keyboard
 - ▶ `./zombies < test1.txt`

OUTPUT REDIRECTION

- ▶ `'>'` takes the stdout (aka would normally go to terminal from cout) and redirects it to a file
- ▶ `./randgen 10 100 > random.txt`

PIPES

- ▶ Using the | (pipe) connects the stdout of one program to the stdin of another
- ▶ Everything you 'cout' in one program becomes the source for 'cin' in the 2nd
- ▶ `./randgen 10 100 | ./average`
- ▶ `wget http://ee.usc.edu/~redekopp/cs103/redirect_pipe.tar`

REDIRECTION AND PIPING EXAMPLES

- ▶ http://ee.usc.edu/~redekopp/cs103/redirect_pipe.tar
- ▶ randgen
- ▶ average

LAB 7 OVERVIEW

- ▶ Lab 7: re-do word scramble game to read the word bank words from a file.
- ▶ File format:
 - ▶ First line: number of words (N)
 - ▶ N words separated by white space



LAB 7 OVERVIEW

- ▶ Attempt to open the file. If that succeeds (`ifile.fail() == false`) go on
- ▶ Read in number of words to a variable
- ▶ If that succeeds (`ifile.fail() == false`), allocate (with dynamic memory) a word back array. `wordBank` will hold `char*`, so it has type `char**`

wordbank.txt


3
computer
trojan hello



wordBank[0]

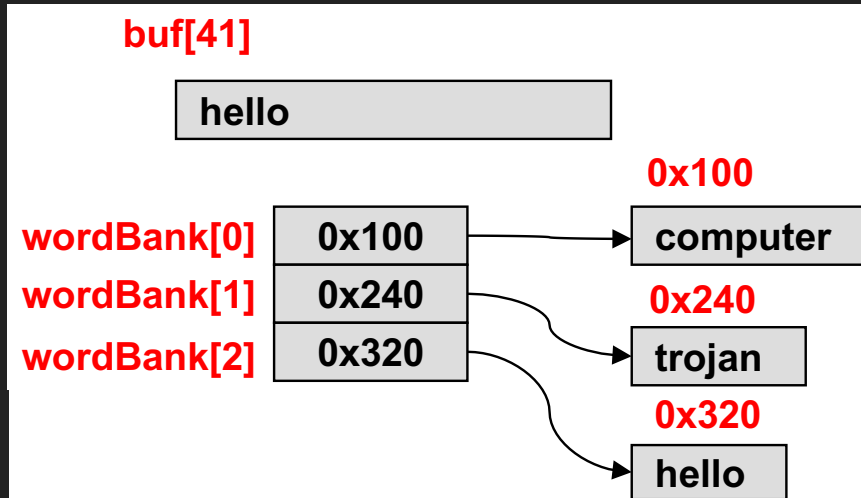
wordBank[1]

wordBank[2]



LAB 7 OVERVIEW

- ▶ In a loop (N times) read in a word to a temporary buffer and then allocate some memory to hold that word.
- ▶ Why two steps? (read then allocate?)
- ▶ Then copy word from temporary buffer to allocated memory



LAB 7 OVERVIEW

- Now you have a dynamically allocated `wordBank[]` that could hold an arbitrary number of words

wordbank.txt

3
computer
trojan hello

